

“Scales are functions that map from an input domain to an output range.”

That’s Mike Bostock’s definition of D3 scales from an earlier version of the D3 docs, and there is no clearer way to say it.

The values in any dataset are unlikely to correspond exactly to pixel measurements for use in your visualization. Scales provide a convenient way to map those data values to new values useful for visualization purposes.

D3 scales are *functions* with parameters that you define. Once they are created, you call the `scale` function and pass it a data value, and it nicely returns a scaled output value. You can define and use as many scales as you like.

It might be tempting to think of a scale as something that appears visually in the final image—like a set of tick marks, indicating a progression of values. *Do not be fooled!* Those tick marks are part of an *axis*, which is a *visual representation* of a scale. A scale is a mathematical relationship, with no direct visual output. I encourage you to think of scales and axes as two different, yet related, elements.

This chapter addresses primarily *linear* scales, because they are most common and easiest understood. Once you understand linear scales, the others—ordinal, logarithmic, square root, and so on—will be a piece of cake.

## Apples and Pixels

Imagine that the following dataset represents the number of apples sold at a roadside fruit stand each month:

```
var dataset = [ 100, 200, 300, 400, 500 ];
```

First of all, this is great news, as the stand is selling 100 additional apples each month! Business is booming. To showcase this success, you want to make a bar chart illustrating the steep upward climb of apple sales, with each data value corresponding to the height of one bar.

Until now, we've used data values directly as display values, ignoring unit differences. So if 500 apples were sold, the corresponding bar would be 500 pixels tall.

That could work, but what about next month, when 600 apples are sold? And a year later, when 1,800 apples are sold? Your audience would have to purchase ever-larger displays just to be able to see the full height of those very tall apple bars! (Mmm, apple bars!)

This is where scales come in. Because apples are not pixels (which are also not oranges), we need scales to translate between them.

## Domains and Ranges

A scale's *input domain* is the range of possible input data values. Given the preceding apple data, appropriate input domains would be either 100 and 500 (the minimum and maximum values of the dataset) or 0 and 500.

A scale's *output range* is the range of possible output values, commonly used as display values in pixel units. The output range is completely up to you, as the information designer. If you decide the shortest apple bar will be 10 pixels tall, and the tallest will be 350 pixels tall, then you could set an output range of 10 and 350.

For example, create a scale with an input domain of  $[100, 500]$  and an output range of  $[10, 350]$ . If you handed the low input value of 100 to that scale, it would return its lowest range value, 10. If you gave it 500, it would spit back 350. If you gave it 300, it would hand 180 back to you on a silver platter. (300 is in the center of the domain, and 180 is in the center of the range.)

We can visualize the domain and range as corresponding axes, side-by-side, displayed in Figure 7-1.

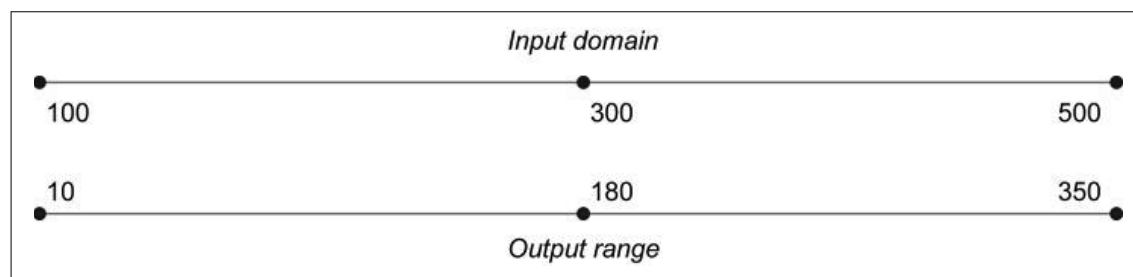


Figure 7-1. An input domain and an output range, visualized as parallel axes

One more thing: to prevent your brain from mixing up the *input domain* and *output range* terminology, I'd like to propose a little exercise. When I say “input,” you say “domain.” Then I say “output,” and you say “range.” Ready? Okay:

Input! *Domain!*

Output! *Range!*

Input! *Domain!*

Output! *Range!*

Got it? Great.

## Normalization

If you're familiar with the concept of *normalization*, it might be helpful to know that, with a linear scale, that's all that is really going on here.

Normalization is the process of mapping a numeric value to a new value between 0 and 1, based on the possible minimum and maximum values. For example, with 365 days in the year, day number 310 maps to about 0.85, or 85% of the way through the year.

With linear scales, we are just letting D3 handle the math of the normalization process. The input value is normalized according to the domain, and then the normalized value is scaled to the output range.

## Creating a Scale

D3's linear scale function generator is accessed with `d3.scaleLinear()`. I recommend opening the sample code page *01\_scale\_test.html* and typing each of the following into the console:

```
var scale = d3.scaleLinear();
```

Congratulations! Now `scale` is a function to which you can pass input values. (Don't be misled by the `var`. Remember that in JavaScript, variables can store functions.)

```
scale(2.5); //Returns 2.5
```

Because we haven't set a domain and a range yet, this function will map input to output on a 1:1 scale. That is, whatever we input will be returned unchanged.

We can set the scale's input domain to 100,500 by passing those values to the `domain()` method as an array. Note the hard brackets indicating an array:

```
scale.domain([100, 500]);
```

Set the output range in similar fashion, with `range()`:

```
scale.range([10, 350]);
```

These steps can be done separately, as just shown, or chained together into one line of code:

```
var scale = d3.scaleLinear()
    .domain([100, 500])
    .range([10, 350]);
```

Either way, our scale is ready to use!

```
scale(100); //Returns 10
scale(300); //Returns 180
scale(500); //Returns 350
```

Typically, you will call scale functions from within an `attr()` method or similar, not on their own. Let's modify our scatterplot visualization to use dynamic scales.

## Scaling the Scatterplot

To revisit our dataset from the scatterplot:

```
var dataset = [
    [5, 20], [480, 90], [250, 50], [100, 33], [330, 95],
    [410, 12], [475, 44], [25, 67], [85, 21], [220, 88]
];
```

You'll recall that this `dataset` is an array of arrays. We mapped the first value in each array onto the x-axis, and the second value onto the y-axis. Let's start with the x-axis.

Just by eyeballing the x values, it looks like they range from 5 to 480, so a reasonable input domain to specify might be 0,500, right?

Why are you giving me that look? Oh, because you want to keep your code flexible and scalable, so it will continue to work even if the data changes in the future. Very smart! Remember, if we were building a data dashboard for the roadside apple stand, we'd want our code to accommodate the enormous projected growth in apple sales. Our chart should work just as well with 5 apples sold as 5 million.

### d3.min() and d3.max()

Instead of specifying fixed values for the domain, we can use the convenient array functions `d3.min()` and `d3.max()` to analyze our dataset on the fly. For example, this loops through each of the x values in our arrays and returns the value of the greatest one:

```
d3.max(dataset, function(d) {
    return d[0]; //References first value in each subarray
});
```

That code will return the value 480, because 480 is the largest x value in our dataset. Let me explain how it works.

Both `min()` and `max()` work the same way, and they can take either one or two arguments. The first argument must be a reference to the array of values you want evaluated, which is `dataset`, in this case. If you have a simple, one-dimensional array of numeric values, like `[7, 8, 4, 5, 2]`, then it's obvious how to compare the values against each other, and no second argument is needed. For example:

```
var simpleDataset = [7, 8, 4, 5, 2];
d3.max(simpleDataset); // Returns 8
```

The `max()` function simply loops through each value in the array, and identifies the largest one.

But our dataset is not just an array of numbers; it is an array of arrays. Calling `d3.max(dataset)` might produce unexpected results:

```
var dataset = [
    [5, 20], [480, 90], [250, 50], [100, 33], [330, 95],
    [410, 12], [475, 44], [25, 67], [85, 21], [220, 88]
];
d3.max(dataset); // Returns [85, 21]. What???
```

To tell `max()` which *specific* values we want compared, we must include a second argument, an *accessor function*:

```
d3.max(dataset, function(d) {
    return d[0];
});
```

The accessor function is an anonymous function to which `max()` hands off each value in the data array, one at a time, as `d`. The accessor function specifies *how to access* the value to be used for the comparison. In this case, our data array is `dataset`, and we want to compare only the `x` values, which are the first values in each subarray, meaning in position `[0]`. So our accessor function looks like this:

```
function(d) {
    return d[0]; //Return the first value in each subarray
}
```

Note that this looks suspiciously similar to the syntax we used when generating our scatterplot circles, which also used anonymous functions to retrieve and return values:

```
.attr("cx", function(d) {
    return d[0];
})
.attr("cy", function(d) {
    return d[1];
})
```

This is a common D3 pattern. Soon you will be very comfortable with all manner of anonymous functions crawling all over your code.

## Setting Up Dynamic Scales

Putting together what we've covered, let's create the scale function for our x-axis:

```
var xScale = d3.scaleLinear()
    .domain([0, d3.max(dataset, function(d) { return d[0]; })])
    .range([0, w]);
```

First, notice that I named it `xScale`. Of course, you can name your scales whatever you want, but a name like `xScale` helps me remember what this function does.

Second, notice that both the domain and range are specified as two-value arrays in hard brackets.

Third, notice that I set the low end of the input domain to 0. (Alternatively, you could use `min()` to calculate a dynamic value.) The upper end of the domain is set to the maximum value in `dataset` (which is currently 480, but could change in the future).

Finally, observe that the output range is set to 0 and `w`, the SVG's width.

We'll use very similar code to create the scale function for the y-axis:

```
var yScale = d3.scaleLinear()
    .domain([0, d3.max(dataset, function(d) { return d[1]; })])
    .range([0, h]);
```

Note that the `max()` function here references `d[1]`, the y value of each subarray. Also, the upper end of `range()` is set to `h` instead of `w`.

The scale functions are in place! Now all we need to do is use them.

## Incorporating Scaled Values

Revisiting our scatterplot code, we now simply modify the original line where we created a circle for each data value. As a reminder, what follows is operating on a selection of newly created circles. These lines set those circles' positions along the x-axis:

```
.attr("cx", function(d) {
    return d[0]; //Returns original value bound from dataset
})
```

to return a scaled value (instead of the original value):

```
.attr("cx", function(d) {
    return xScale(d[0]); //Returns scaled value
})
```

Likewise, for the y-axis, this:

```
.attr("cy", function(d) {
    return d[1];
})
```

is modified as:

```
.attr("cy", function(d) {  
    return yScale(d[1]);  
})
```

For good measure, let's make the same change where we set the coordinates for the text labels, so these lines:

```
.attr("x", function(d) {  
    return d[0];  
})  
.attr("y", function(d) {  
    return d[1];  
})
```

become this:

```
.attr("x", function(d) {  
    return xScale(d[0]);  
})  
.attr("y", function(d) {  
    return yScale(d[1]);  
})
```

And there we are!

Check out the working code in *02\_scaled\_plot.html*. Visually, the result in Figure 7-2 is disappointingly similar to our original scatterplot! Yet we are making more progress than might be apparent.




Figure 7-2. Scatterplot using *x* and *y* scales

## Refining the Plot

You might have noticed that smaller *y* values are at the top of the plot, and the larger *y* values are toward the bottom. Now that we're using D3 scales, it's super easy to reverse that, so greater values are higher up, as you would expect. It's just a matter of changing the output range of *yScale* from:

```
.range([0, h]);
```

to:

```
.range([h, 0]);
```

See *03\_scaled\_plot\_inverted.html* for the code that results in Figure 7-3.




Figure 7-3. Scatterplot with *y* scale inverted

Yes, now a *smaller* input to `yScale` will produce a *larger* output value, thereby pushing those circles and text elements down, closer to the base of the image. I know, it's almost too easy!

Yet some elements are getting cut off. Let's introduce a `padding` variable:

```
var padding = 20;
```

Then we'll incorporate the `padding` amount when setting the range of both scales. This will help push our elements in, away from the edges of the SVG, to prevent them from being clipped.

The range for `xScale` was `range([0, w])`, but now it's:

```
.range([padding, w - padding]);
```

The range for `yScale` was `range([h, 0])`, but now it's:

```
.range([h - padding, padding]);
```

This should provide us with 20 pixels of extra room on the left, right, top, and bottom edges of the SVG. And it does; see Figure 7-4.




Figure 7-4. Scatterplot with *padding*

But the text labels on the far right are still getting cut off, so I'll double the amount of `xScale`'s padding on the right side by multiplying by two to achieve the result shown in Figure 7-5.

```
.range([padding, w - padding * 2]);
```




Figure 7-5. Scatterplot with *more padding*





The way I've introduced padding here is simple, but not elegant. Eventually, you'll want more control over how much padding is on each side of your charts (top, right, bottom, left), and it's useful to standardize how you specify those values across projects. Although I haven't used Mike Bostock's margin convention for the code samples in this book, I recommend taking a look to see if it could work for you.

Better! Reference the code so far in *04\_scaled\_plot\_padding.html*. But there's one more change I'd like to make. Instead of setting the radius of each circle as the square root of its y value (which was a bit of a hack, and not visually useful), why not create another custom scale?

```
var rScale = d3.scaleLinear()
    .domain([0, d3.max(dataset, function(d) { return d[1]; })])
    .range([2, 5]);
```

Now we can use `rScale` to scale each circle's radius. (My admonishment to always scale circles by *area* and not *radius* still stands. Bear with me for a moment; we'll come back to this.)

```
.attr("r", function(d) {
    return rScale(d[1]);
});
```

This is exciting because we are guaranteeing that our radius values will *always* fall within the range of 2,5. (Or *almost* always; see reference to `clamp()` later.) So data values of 0 (the minimum input) will get circles of radius 2 (or a diameter of 4 pixels). The very largest data value will get a circle of radius 5 (diameter of 10 pixels).

Voila: Figure 7-6 shows our first scale used for a visual property other than an axis value. (See *05\_scaled\_plot\_radii.html*.)




Figure 7-6. Scatterplot with scaled radii

Finally, just in case the power of scales hasn't yet blown your mind, I'd like to add one more array to the dataset: `[600, 150]`.

Boom! Check out *06\_scaled\_plot\_big.html*. Notice how all the old points in Figure 7-7 maintained their relative positions but have migrated closer together, down and to the left, to accommodate the newcomer in the top-right corner.




Figure 7-7. Scatterplot with big numbers added

And now, one final revelation: we can now very easily change the size of our SVG, and *everything scales accordingly*. In Figure 7-8, I've increased the value of *h* from 100 to 300 and made *no other changes*.



Figure 7-8. Large, scaled scatterplot

Boom, again! See [07\\_scaled\\_plot\\_large.html](#). Hopefully, you are seeing this and realizing: no more late nights tweaking your code because the client decided the graphic should be 800 pixels wide instead of 600. Yes, you will get more sleep because of me (and D3's brilliant built-in methods). Being well-rested is a competitive advantage. You can thank me later.

## Other Methods

`d3.scaleLinear()` has several other handy methods that deserve a brief mention here:

### `nice()`

This tells the scale to take whatever input domain that you gave to `domain()` and expand both ends to the nearest round value. From the D3 wiki: “For example, for a domain of `[0.201479..., 0.996679...]`, a nice domain might be `[0.2, 1.0]`.” This is useful for humans, who are not computers and find it hard to read numbers like `0.20147987687960267`.

### `rangeRound()`

Use `rangeRound()` in place of `range()`, and all values output by the scale will be rounded to the nearest whole number. This is useful if you want shapes to have exact pixel values, to avoid the fuzzy edges that could arise with antialiasing.

### `clamp()`

By default, a linear scale *can* return values outside of the specified range. For example, if given a value outside of its expected input domain, a scale will return a number also outside of the output range. Calling `clamp(true)` on a scale, however, forces all output values to be within the specified range. This means excessive values will be rounded to the range’s low or high value (whichever is nearest).

To use any of these special methods, just tack them onto the chain in which you define the original scale function. For example, to use `nice()`:

```
var scale = d3.scaleLinear()
    .domain([0.123, 4.567])
    .range([0, 500])
    .nice();
```

## Other Scales

In addition to linear scales (discussed earlier), D3 has several other built-in scale methods:

### `scaleSqrt`

A square root scale.

### `scalePow`

A power scale (good for the gym, er, I mean, useful when working with exponential series of values, as in “to the power of” some exponent).

`scaleLog`

A logarithmic scale.

`scaleQuantize`

A linear scale with discrete values for its output range, for when you want to sort data into “buckets.”

`scaleQuantile`

Similar to `scaleQuantize`, but with discrete values for its input domain (when you already have “buckets”).

`scaleOrdinal`

Ordinal scales use nonquantitative values (like category names) for output; perfect for comparing apples and oranges.

`schemeCategory10`, `schemeCategory20`, `schemeCategory20b`, and

`schemeCategory20c`

Handy preset ordinal scales that output either 10 or 20 categorical colors.

`scaleTime`

A scale method for date and time values, with special handling of ticks for dates.

Before we move on to axes, let’s explore how to use two of these: square root and time scales.

## Square Root Scales

In Chapter 6, I mentioned how circles should always be scaled by *area*, not by *radius* value, to better match how we perceive circle sizes. To get from data value (effectively the “area” value) to a radius value, we used `Math.sqrt()`, like so:

```
.attr("r", function(d) {  
  return Math.sqrt(d);  
});
```

If we used a square root scale instead, the code would look a little different:

```
.attr("r", function(d) {  
  return aScale(d); // 'a' scale for 'area'  
});
```

The main benefit is we get to take advantage of the scale’s ability to set up a domain and a range (and to update those in the future, should our needs change). Why not let the scale handle all the math?

Let’s replace our radius `rScale` with an area `aScale` as follows:

```
var aScale = d3.scaleSqrt() // <--New!  
  .domain([0, d3.max(dataset, function(d) { return d[1]; })))  
  .range([0, 10]); // <--New!
```

There are only three changes here:

1. Switching `scaleLinear` to `scaleSqrt`.
2. Renaming the scale to `aScale`. This is still being used to set each circle's `r` radius value (because there is no `circle-area` attribute in SVG), but the value is calculated with the data value mapping to *area*, per best practices.
3. Adjusting the `range`, which I only did to make the relative differences in size more obvious. Note that the values 0 and 10 are arbitrary. Remember, what matters is the *relative* size (areas) of the circles, not the actual or absolute sizes (areas).

See [08\\_scaled\\_plot\\_sqrt\\_scale.html](#) for the result, which looks like Figure 7-9.




Figure 7-9. Using a square root scale for circle areas

## Time Scales

We frequently need to visualize data with a temporal dimension. Fortunately, D3 has you covered, although there are a few extra steps required.

First, meet the JavaScript `Date` object. It's how JavaScript understands time. Open your console and type `new Date`. In return, you'll see a string representation of a new `Date` object that represents whatever date and time it is *right now*.

```
Mon Feb 20 2017 13:44:00 GMT-0800 (PST)
```

As a smart human being, you can mentally parse this to understand I was writing this chapter on February 20, 2017, at 1:44 pm. That happened to be a Monday, and I'm in the Pacific Standard Time zone, which is 8 hours behind Coordinated Universal Time.

(I won't go down the temporal rabbit hole here, but it can be enlightening to read about UTC, the nature of time, and various attempts to corral this relative concept in code, such as the wonderful *Moment.js*. I also highly recommend episode #70 of the *Data Stories* podcast with Enrico Bertini and Moritz Stefaner, "Rocket Science with Rachel Binx," in which Rachel discusses the challenges of coordinating interplanetary time measurements.)

JavaScript and D3 can only perform time and date calculations on *Date* objects, not on strings (even strings that look very much like dates, to human eyes). So working with dates in D3 involves:

1. Converting strings to *Date* objects
2. Using time scales, as needed
3. Formatting *Date* objects as human-friendly strings, for display to the user

Let's take these one at a time.

### Converting strings to dates

I've provided some dummy data as a time series in *time\_scale\_data.csv*, which looks like this:

```
Date,Amount
01/01/17,35
01/02/17,30
01/03/17,24
01/04/17,37
01/05/17,54
...
```

Each row represents a daily measurement, with one *Amount* for each day in January 2017. (The data may be random and meaningless, but at least it's timely!)

As a human who can interpret context, you've already deduced that these dates are in the format month/day/year, a sequence common in the United States and pretty much nowhere else. In order to tell the context-insensitive computer how to interpret these strings, we write:

```
//For converting strings to Dates
var parseTime = d3.timeParse("%m/%d/%y");
```

The strange syntax of "%m/%d/%y" tells D3 to look for three values, separated by slashes: month with leading zero, day of the month with leading zero, and two-digit year number.

Personally, I can never remember all the various signifiers for each date/time component. I recommend the API reference for all possible time formatting values.

To verify this really works, try opening *09\_time\_scale.html* in your browser. Then, in the console, try parsing your birthday into a `Date` object by using the format `mm/dd/yy`. For example:

```
parseTime("02/20/17")  
//Returns: Mon Feb 20 2017 00:00:00 GMT-0800 (PST)
```

Since we specified only a two-digit year, JavaScript has to guess about which century we intended: the 21st, or the 20th? If you were born in 1969 or since then, JavaScript will guess right:

```
parseTime("02/20/69")  
//Returns: Thu Feb 20 1969 00:00:00 GMT-0800 (PST)
```

If you were born in 1968 or earlier, however, it may guess you are from the future:

```
parseTime("02/20/68")  
//Returns: Mon Feb 20 2068 00:00:00 GMT-0800 (PST)
```

Finally, if you really *are* from the future, please email me. I have lots of questions.

In all seriousness, to avoid ambiguity with dates, I strongly recommend using four-digit years and `%Y`. Use your spreadsheet to reformat date values before exporting to CSV for D3.

Using a technique described in Chapter 5, I'll define a row conversion function. Once our CSV data is loaded in, this function is called by `d3.csv()` on each row. We have just two values per row (date and amount), and this specifies how to convert each from a string to a `Date` object or number, respectively. (It's easier to do this conversion now, rather than remember to do it later.) Note how we pass the string `d.Date` to `parseTime()`, and it returns a `Date` object.

```
//Function for converting CSV values from strings to Dates and numbers  
var rowConverter = function(d) {  
  return {  
    Date: parseTime(d.Date),  
    Amount: parseInt(d.Amount)  
  };  
}
```

To verify that this worked, I jump to the console and type `dataset`. See the result in Figure 7-10.

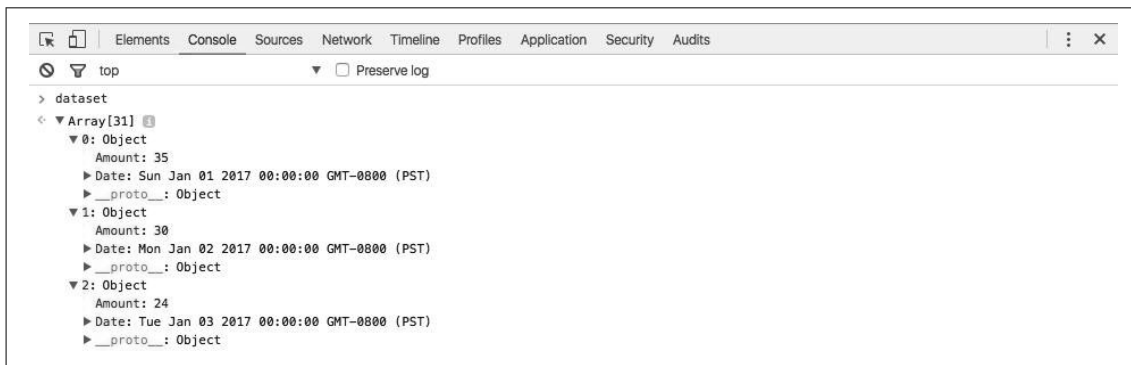


Figure 7-10. Look mom, no strings!

## Scaling time

The hard part is over. Our data is loaded in, and the strings are now dates. We use `scaleTime()` to define a time scale:

```
xScale = d3.scaleTime()
    .domain([
        d3.min(dataset, function(d) { return d.Date; }),
        d3.max(dataset, function(d) { return d.Date; })
    ])
    .range([padding, w - padding]);
```

You'll recognize `d3.min()` and `d3.max()` from earlier. How cool that they work on dates, and not just simple, numeric values! We can verify that the domain matches our dataset by typing `xScale.domain()` in the console, as in Figure 7-11.




Figure 7-11. Verifying the `xScale` domain runs from January 1 through January 31, 2017

Finally, as with any scale, there is one more step: to actually call the scale, to convert data values as needed. In this case, I'm using:

```
.attr("cx", function(d) {
    return xScale(d.Date);
})
```

The only change here is that we're specifying `d.Date` as the value to be scaled. See the final chart in Figure 7-12.






Figure 7-12. Time-scaled circles

It's not clean or beautiful (or meaningful), but it works! You can see the circles being positioned in chronological order, from left to right, and also positioned vertically, per their Amount values.

### Formatting dates and strings

To create the labels on that chart, we need to convert our Date values back to human-readable strings. I've chosen to use the format of the month name's three-letter abbreviation followed by the day of the month. This is defined in a new time formatting function as:

```
//For converting Dates to strings
var formatTime = d3.timeFormat("%b %e");
```

Again, please see the API reference for time formatting options. You can take my word for it that %b results in Jan, Feb, Mar, and so on, while %e results in 1, 2, 3, and so on up until the end of the month.

Having specified the formatter, we call it whenever we need those human-readable strings generated, such as for inserting the text for those labels:

```
.text(function(d) {
    return formatTime(d.Date);
})
```

Explore the working code in [09\\_time\\_scale.html](#). Try tweaking the time formatter options. Can you get it to return the full month name? How about the four-digit year? How about the time zone or day of the week? Once you are comfortable with time scales, you'll see it's no harder working on a small scale (seconds) than on a large one (centuries).

Now that you have mastered the power of scales, it's time to express them visually as, yes, *axes*!